


Pooled human preference data

Per-annotator-calibrated reward



At deployment, annotator j 's predicted rating on item x : $\hat{y}_{ji} = \hat{\alpha}_j^{\text{shrunk}} \cdot r_\phi(x) + \hat{\beta}_j^{\text{shrunk}}$